

Revision Report: Response to Comments on the Manuscript

Rashedul Arefin Ifty

August 2026

1 Summary

This report records the revisions made to the manuscript following the written advice received in August 2026. The advice identified five recurring problems in the presentation, namely inconsistent word choice, sentences carrying more than one message, sentences too short to be understood, ambiguous pronouns, and unclear subsection structure. Fifteen specific examples were named. All fifteen have been addressed. The revision also covers further occurrences of the same patterns that were found by searching the whole manuscript.

The manuscript builds without errors and remains within ten pages. No numerical result, table entry, or claim was changed by this revision.

2 Word Choice

The advice asked that unnecessary technical terms be removed, and that any term which is kept be defined once and then used consistently. The following changes were made.

Term	Action
signals, pipeline interfaces, cascade channel	Removed. The sentence in Section II-A was rewritten to explain the connection between the models in plain language.
terminal	Removed from 15 places. Replaced by “planning” where the last model is meant, and by “the last model” where the general pipeline is meant.
tolerance	Replaced by “threshold” in 4 places, including the input line of Algorithm 1.
end-to-end literature	Rewritten as “earlier studies on end-to-end driving models”.
one fully controlled update	Replaced by “a single update”.
injected effect	Removed from 8 places. Replaced by “the change caused by the retraining”.
drift score	Defined at first use in Section III-B, and now used consistently. Previously the same quantity was called the interface-drift score, the drift score, the front-end, the drift screen, and the drift proxy.
profile	Removed from 10 places. Replaced by “the two-mode measurement”. The term was never defined, and a cross-reference pointed to a section that no longer contained it.

Term	Action
effect vector	Removed. Replaced by “the change at every downstream model”.
regime	Removed from Sections III to V, including a table column heading.
frozen	Removed from 10 places, matching the plain wording adopted in Sections I and II. Replaced by “not retrained” or “unchanged”.
horizon	Removed from Section IV onward. Replaced by an explicit number of seconds.
interface	Removed from 18 places in our own description of the method. The term is retained only where it describes the work of other authors in the related-work section.

A further consistency pass replaced “driving pipeline” and “driving output” with “AV pipeline” and “planning output”, so that one name is used for each concept across the whole manuscript.

3 Sentences Carrying More Than One Message

The four sentences named in the advice were each split into shorter sentences, with one point per sentence.

1. The last sentence of Section II-A, on how a change reaches a model that was not retrained, is now three sentences.
2. The last sentence of Section IV-C, on the L2 metric and its caveat, is now three sentences that separate what the protocol is, that it has known limitations, and how it is used here.
3. The sentence in Section V-E on the predictability of the outcome is now two sentences.
4. The sentence in Section V-G on the AUROC of the drift score is now two sentences, and the value is stated together with what it means.
5. An additional sentence in Section III-A, on why all the changes are reported together, was split in the same way.

4 Sentences That Lacked Information

The four sentences named in the advice were expanded.

1. **Fine-tuning data size.** The description now states what the half and the quarter subsets are, how they are selected, that the same subsets are used in every run, and why this factor is varied.
2. **The isolation check.** The claim that the check is exact is now supported by the reasoning, namely that both models are deterministic and have no trainable parameters, so identical inputs give identical outputs.
3. **Figure reference.** The sentence “Fig. 2 shows the signature” was replaced. The values are now stated first and the figure is referenced afterwards.

4. **Divergence of the two metrics.** The sentence now states which metric moved in which direction.

5 Ambiguous Pronouns

The three sentences named in the advice were rewritten so that every reference names the thing it refers to.

1. The last sentence of Section II-B, which contained “the two” and “it”, now names entangled enhancement and ordinary regression explicitly.
2. The sentence in Section II-C beginning “It is what lets us...” now begins with “Isolated mode”.
3. The sentence in Section V-H, which used “it” for two different things, was rewritten as two sentences that name CARA and the disagreement between the two metrics.

6 Subsection Structure

The two subsection titles named in the advice were changed so that each title states what the subsection contains.

Before	After
Interface-Drift Front-End	Low-Cost Drift Screen
Why the Method Needs a Pipeline Architecture	Applicability to Pipeline Architectures

Two further structural problems were corrected. The overview of Section V stated that seven experiments follow, while eight are described; the count was corrected. Section IV-E repeated material already given in Sections IV-B and IV-D, and was reduced to the part that is not stated elsewhere.

7 Other Corrections

- The role of the drift score is now stated explicitly. It is a screening statistic and not the admission criterion, and the admission decision uses the change at the last model.
- Grammatical corrections were made throughout the abstract, the introduction, and Section II, including subject-verb agreement, article use, and the consistent spelling of “entangled enhancement”.
- Redundant passages were removed from Sections V-H, V-I, VI-B, and the conclusion, where the same point was made twice.

8 Work Remaining

The advice noted that the named examples were samples and that further sentences require revision. The following are known and not yet addressed.

- Approximately thirteen further sentences show the same patterns of packed content, terseness, or ambiguous reference, in Sections V, VI, VII, and VIII.
- Section III-C remains an argument about applicability placed among subsections that describe the method itself. Only its title has been changed.
- Section IV-A covers the dataset, the splits, the audit set, the labels, and all three models under a single heading.
- Two flagged terms remain in Section II, which was revised by hand and left unchanged in this pass.